




All methods: two causal observations, DDPM100, 16-action prediction / 8-action execution.
One training replicate; model count and total compute differ across methods.



Framework / data



API-authored decisions



Learned policy